

Simulator: Tensor-ALU Module

```
1 #procedure input: #balls (n); (kx, ky); ball radius (r)
2 device kernel():
3     i,j=tile co-ordinate based on thread ids
4     if not_halo:
5         worklist <- worklist_new #stage1
6         balls <- balls_new
7         for b in worklist(i,j):
8             for n in eight_neighbouring_tiles of (i,j):
9                 if overlap(n,b):
10                     if b is 0:
11                         die()
12                     naive_collider(n,b) #updates position, velocity of 'balls'
13 if first_thread_of_a_block:
14     atomicInc(*counter,max_counter); #init: counter=0, max_counter=total # blocks-1
15 __syncthreads(); while(*counter); #global barrier
16 for b in worklist(i,j): #stage2
17     balls_new.position <- balls.position + (p==0)? keyboard_moves:balls.velocity
18     balls_new.velocity <- balls.velocity
19     reflect 'b' if it is at/outside boundary
20     i_new, j_new <- new_tile_of_ball(b)
21     atomically: worklist_new[i_new,j_new] = b
22
23 main():
24     balls[0:n] <- initBalls() #randomized positions, velocities
25     glutInit(); glewInit() #display, keyboard, timer, view-port inits
26     while(True):
27         if window_is_reshaped:
28             tiles <- split window into rectangles based on (kx, ky, r)
29             worklist <- assign balls to tiles based on their position
30             device kernel(tiles assigned to GPU cores from threads->blocks in 2D) in stream1
31             if approp. keyboard_input: convolution_kernel() in stream2 to modify background
32             glutDraw(balls)
```

- Worklist
- Shared variables
- Structure of Arrays
- Pinned memory (`cudaHostAlloc`)
- Global barrier using voting and `sync_threads()`
- `nvprof`, `nsys`,
- Texture memory
- lock-step execution ternary operator